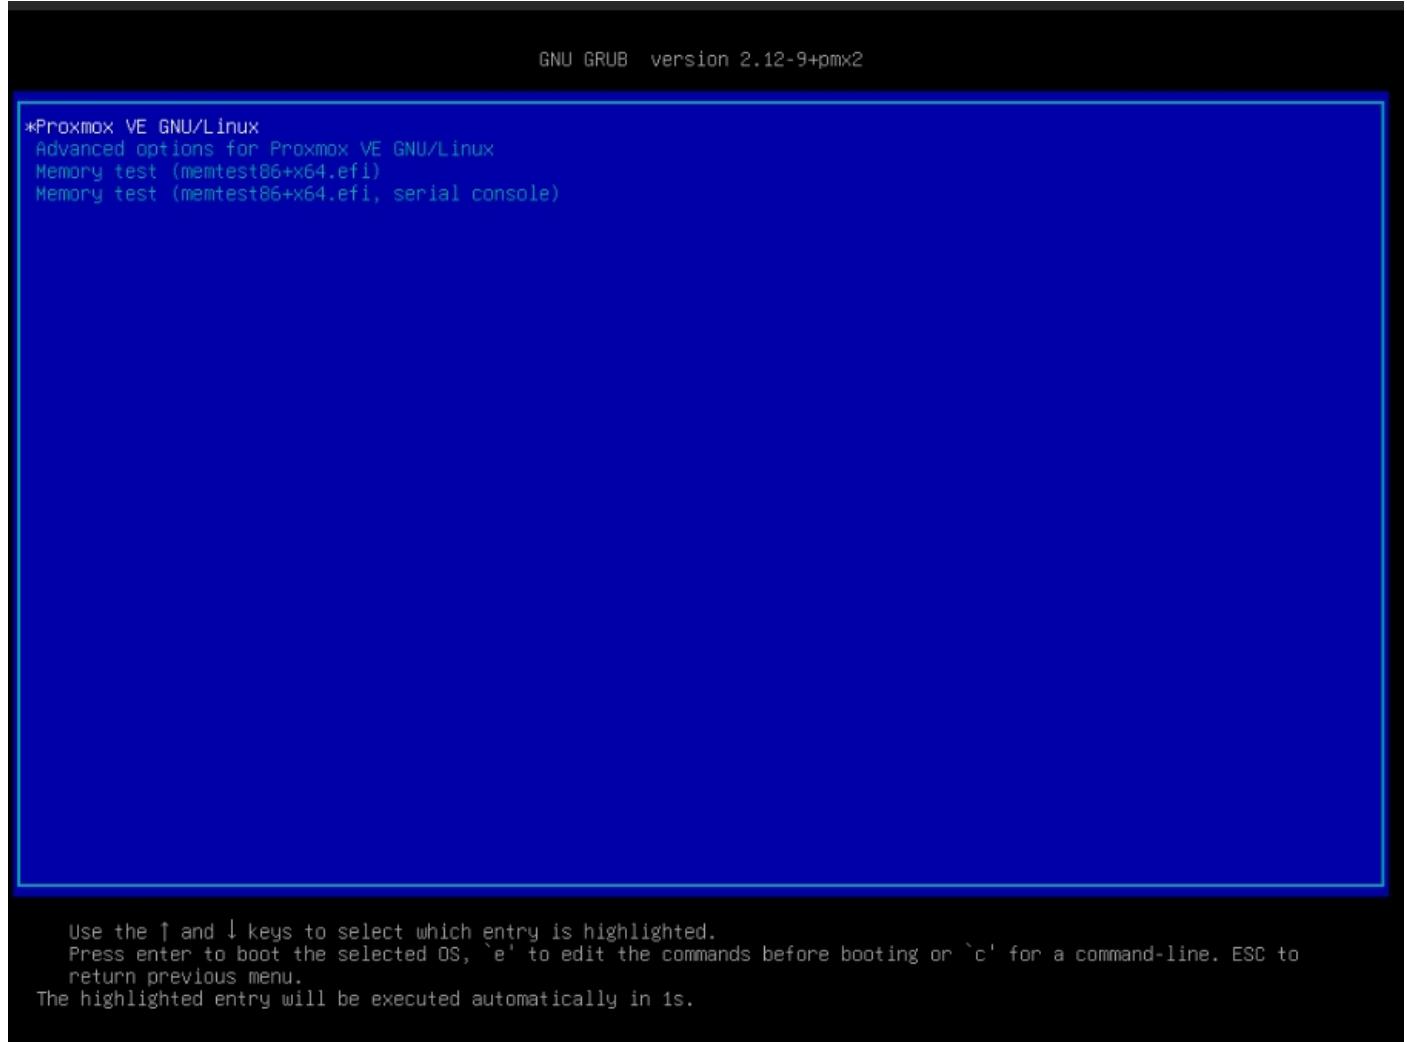


Proxmox + CT ollama + VM ubuntu docker OpenwebUI



After proxmov deployment

- 4 lspci | grep NVIDIA
- 5 lsmod | grep nouveau ### nouveau 要禁用
- 6 cd /
- 7 ls -l /dev/dri
- 8 ls -l /dev/nvidia*
- 9 echo -e "blacklist nouveau\noptions nouveau modeset=0" > /etc/modprobe.d/blacklist-nouveau.conf && update-initramfs -u && reboot
- 10 cat /etc/modprobe.d/blacklist-nouveau.conf

```
root@pve:~# cat /etc/modprobe.d/blacklist-nouveau.conf
blacklist nouveau
```

```
options nouveau modeset=0
```

```
11 lsmod | grep nouveau
```

```
12 lsmod | grep nouveau
```

```
13 lspci | grep NVIDIA
```

```
wget https://us.download.nvidia.com/XFree86/Linux-x86\_64/580.95.05/NVIDIA-Linux-x86\_64-580.95.05.run
```

```
121 chmod +x NVIDIA-Linux-x86_64-580.95.05.run
```

```
123 ./NVIDIA-Linux-x86_64-580.95.05.run
```

```
124 reboot
```

```
Nvidia-smi verify driver version
```

```
###PVE 完成
```

```
## create CT
```

```
61 pveam download local ubuntu-24.04-standard_24.04-2_amd64.tar.zst
```

```
Create CT101 ollama
```

```
root@pve:~# cat /etc/pve/lxc/101.conf
```

```
# === [ NVIDIA GPU %E6%98%A0%E5%B0%84%E8%A8%AD%E5%AE%9A -  
%E8%AB%8B%E5%8B%99%E5%BF%85%E8%B2%BC%E5%9C%A8%E9%85%8D%E7%BD%AE%E6  
%AA%94%E6%9C%AB%E5%B0%BE ] ===
```

```
# %E6%AC%8A%E9%99%90%E6%8E%88%E6%AC%8A  
(%E4%BD%BF%E7%94%A8%E6%82%A8%E7%A2%BA%E8%AA%8D%E9%81%8E%E7%9A%84%E7  
%B7%A8%E8%99%9F)
```

```
# %E6%98%A0%E5%B0%84 NVIDIA %E5%B0%88%E7%94%A8%E8%A3%9D%E7%BD%AE
```

```
# %E6%98%A0%E5%B0%84 DRM %E6%B8%B2%E6%9F%93%E8%A3%9D%E7%BD%AE
```

```
arch: amd64
```

```
cores: 16
```

```
features: nesting=1
```

```
hostname: ollama-ct
```

```
memory: 32768
```

```
nameserver: 192.168.55.254
```

```
net0:
```

```
name=eth0,bridge=vmbr0,gw=192.168.55.254,hwaddr=BC:24:11:BF:DD:D1,ip=192.168.55.181/24,ty  
pe=veth
```

```
ostype: ubuntu
```

```
rootfs: local-lvm:vm-101-disk-0,size=200G
```

```
searchdomain: nx3230.local
```

```
swap: 1024
```

```
unprivileged: 1
```

```
lxc.cgroup2.devices.allow: c 195:* rwm
```

```
lxc.cgroup2.devices.allow: c 511:* rwm  
lxc.cgroup2.devices.allow: c 226:* rwm  
lxc.mount.entry: /dev/nvidia0 dev/nvidia0 none bind,optional,create=file  
lxc.mount.entry: /dev/nvidiactl dev/nvidiactl none bind,optional,create=file  
lxc.mount.entry: /dev/nvidia-uvm dev/nvidia-uvm none bind,optional,create=file  
lxc.mount.entry: /dev/nvidia-uvm-tools dev/nvidia-uvm-tools none bind,optional,create=file  
lxc.mount.entry: /dev/dri/card0 dev/dri/card0 none bind,optional,create=file  
lxc.mount.entry: /dev/dri/card1 dev/dri/card1 none bind,optional,create=file  
lxc.mount.entry: /dev/dri/renderD128 dev/dri/renderD128 none bind,optional,create=file  
root@pve:~#
```

195:

● 範例：映射 NVIDIA GPU

對於 NVIDIA GPU，LXC 容器設定通常會使用一個特殊選項來處理所有必要的節點：

```
lxc.cgroup2.devices.allow: c 195:* rwm # 允許 NVIDIA 設備  
lxc.mount.entry: /dev/nvidia-uvm dev/nvidia-uvm none bind,optional,create=file  
lxc.mount.entry: /dev/nvidia0 dev/nvidia0 none bind,optional,create=file  
lxc.mount.entry: /dev/nvidiactl dev/nvidiactl none bind,optional,create=file
```

或者，您可以使用 Proxmox 的 `lxc.hook` 腳本來自動處理。

511

您的下一個步驟是執行 **步驟 1: 驗證 NVIDIA 專用裝置節點**，然後開始編輯您的 CT 配置檔。

```
root@pve:~# ls -l /dev/nvidia*
crw-rw-rw- 1 root root 195, 0 Dec 2 16:09 /dev/nvidia0
crw-rw-rw- 1 root root 195, 255 Dec 2 16:09
/dev/nvidiactl
crw-rw-rw- 1 root root 511, 0 Dec 2 16:09 /dev/nvidia-...
```

- ◆ 太好了！您提供的輸出確認了所有 NVIDIA 專有裝置節點都已經在 Proxmox 主機上成功創建並準備就緒。

這些節點（尤其是 `/dev/nvidia0`、`/dev/nvidiactl`、`/dev/nvidia-uvm`）是容器 (CT) 存取 GPU 進行計算或轉碼的關鍵。

現在，我們可以精確地配置您的 CT 映射檔案。

226

```
root@pve:~# ls -l /dev/dri
total 0
drwxr-xr-x 2 root root    ✓      100 Dec 3 13:38 by-path
crw-rw---- 1 root video 226, 0 Dec 3 13:38 card0
crw-rw---- 1 root video 226, 1 Dec 3 13:38 card1
crw-rw---- 1 root render 226, 128 Dec 3 13:38 renderD128
```

c 195:* 和 **c 511:*** 是根據您輸出的 Major Number (例如 195, 0 和 511, 0) 設置的 NVIDIA 專有裝置權限。

c 226:* 是針對您 `/dev/dri/` 下的 顯示渲染裝置 (DRM) 權限。

`lxc.mount.entry` 負責將主機上的實體檔案路徑映射到容器內部相同路徑。

##進到 CT

86 pct start 101

87 pct enter 101

5 vi /etc/apt/sources.list

```
root@ollama-ct:~# cat /etc/apt/sources.list
```

```
deb http://archive.ubuntu.com/ubuntu noble main restricted universe multiverse
```

```
deb http://archive.ubuntu.com/ubuntu noble-updates main restricted universe multiverse
```

```
deb http://archive.ubuntu.com/ubuntu noble-security main restricted universe multiverse
```

```
deb http://archive.ubuntu.com/ubuntu noble main restricted universe multiverse
```

```
deb http://archive.ubuntu.com/ubuntu noble-updates main restricted universe multiverse
```

```
deb http://archive.ubuntu.com/ubuntu noble-security main restricted universe multiverse
```

8 apt install libnvidia-cuda-toolkit

9 apt install libnvidia-cuda-toolkit libnvidia-ml-dev -y

11 apt install nvidia-utils-580 libnvidia-compute-580 libnvidia-decode-580 -y

12 curl -fsSL https://ollama.com/install.sh | sh

13 apt install curl

14 curl -fsSL https://ollama.com/install.sh | sh

15 ollama

34 apt search libnvidia-compute | grep 580

35 apt install nvidia-utils-580

36 nvidia-smi #####CT 與 PVE 版本要一到才跑得起來

37 ollama run llama3

```
root@ollama-ct:/etc/systemd/system# pwd
```

```
/etc/systemd/system
```

```
root@ollama-ct:/etc/systemd/system# cat ollama.service
```

```
[Unit]
```

Description=Ollama Service

After=network-online.target

[Service]

ExecStart=/usr/bin/ollama serve

User=ollama

Group=ollama

Restart=always

RestartSec=3

Environment="PATH=/sbin:/bin:/usr/sbin:/usr/bin"

Environment="OLLAMA_HOST=0.0.0.0"

####這個後面或現在要改 沒有加只有

127.0.0.1 自己能存取，要改 openwebui 才拿得到

[Install]

WantedBy=default.target

root@ollama-ct:/etc/systemd/system#

74 sudo systemctl daemon-reload

75 sudo systemctl restart ollama

##CT ollama 完成

PVE CVM install Ubuntu OS

```
user@vm-103:~$ history
```

```
1 sudo apt update

4 sudo apt install ca-certificates curl gnupg lsb-release -y

5 sudo install -m 0755 -d /etc/apt/keyrings

6 curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o
/etc/apt/keyrings/docker.gpg

7 sudo chmod a+r /etc/apt/keyrings/docker.gpg

8 echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \

9 $(./etc/os-release && echo "$VERSION_CODENAME") stable" | sudo tee
/etc/apt/sources.list.d/docker.list > /dev/null

10 sudo apt update

11 sudo apt install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-
plugin -y

12 sudo docker run hello-world

13 sudo usermod -aG docker $USER

14 docker run -d --name open-webui -p 3000:8080 -e
OLLAMA_BASE_URL=http://192.168.55.181:11434 -v open-webui-data:/app/backend/data
--restart always ghcr.io/open-webui/open-webui:main

15 exit <<<<<跳出重登 !13 才會生效

16 docker ps

17 docker run

18 docker run -d --name open-webui -p 3000:8080 -e
OLLAMA_BASE_URL=http://192.168.55.181:11434 -v open-webui-data:/app/backend/data
--restart always ghcr.io/open-webui/open-webui:main

19 docker ps
```

22 sudo ufw status

Admin Panel • Open WebUI

Not secure 192.168.55.183:3000/admin/settings/connections

Users Evaluations Functions Settings

General

Connections

Models

Evaluations

External Tools

Documents

Web Search

Code Execution

Interface

Audio

Images

Pipelines

Database

General

OpenAI API

Manage OpenAI API Connections

<https://api.openai.com/v1>

Ollama API

Manage Ollama API Connections

<http://192.168.55.181:11434>

Trouble accessing Ollama? [Click here for help.](#)

Direct Connections

Direct Connections allow users to connect to their own OpenAI compatible API endpoints.

Cache Base Model List

Base Model List Cache speeds up access by fetching base models only at startup or on settings save—faster, but may not show recent base model changes.

check 不用改

Direct Connections toggle switch (off)

Cache Base Model List toggle switch (off)

對映到 CT 已經是 Environment="OLLAMA_HOST=0.0.0.0"

如此才能確保 OpenWebUI 能連到 Ollama-IP :11434

下載 Model 的功能才接下去

Not secure https://192.168.55.180:8006/#v10:=qemu%2F103:....

X PROXMOX Virtual Environment 9.1.1

Server View Documentation Create VM Create CT root@pam

Virtual Machine 103 (openwebui-VM) on node 'pve' No Tags

Start Shutdown Console More Help

Datacenter

pve

- 101 (ollama-ct)
- 102 (openwebui-ct)
- 103 (openwebui-VM)**
- localnetwork (pve)
- local (pve)
- local-lvm (pve)

Summary

Console Hardware Cloud-Init Options Task History Monitor Backup Replication Snapshots Firewall Permissions

openwebui-VM (Uptime: 00:07:48)

i Status	running
Heartbeat HA State	none
Node	pve
CPU usage	0.33% of 8 CPU(s)
Memory usage	22.29% (1.78 GiB of 8.00 GiB)
Host memory usage	2.10 GiB
Bootdisk size	100.00 GiB
IPs	No Guest Agent configured

Notes

VM ubuntu docker openwebui

CPU Usage

CPU usage

Hour Maximum Average

Not secure https://192.168.55.180:8006/#v10:=qemu%2F103:....

X PROXMOX Virtual Environment 9.1.1

Server View Documentation Create VM Create CT root@pam

Container 101 (ollama-ct) on node 'pve' No Tags

Start Shutdown Console More Help

Datacenter

pve

- 101 (ollama-ct)**
- 102 (openwebui-ct)
- 103 (openwebui-VM)
- localnetwork (pve)
- local (pve)
- local-lvm (pve)

Summary

Console Resources Network DNS Options Task History Backup Replication Snapshots Firewall Permissions

ollama-ct (Uptime: 00:16:00)

i Status	running
Heartbeat HA State	none
Node	pve
Unprivileged	Yes
CPU usage	0.00% of 16 CPU(s)
Memory usage	0.16% (51.14 MiB of 32.00 GiB)
SWAP usage	0.00% (0 B of 1.00 GiB)
Bootdisk size	11.02% (21.57 GiB of 195.80 GiB)
IPs	192.168.55.181 fe80::be24:11ff:febfd:ddd1

Ubuntu Notes

ollama CT 有對映到GPU

CPU Usage

CPU usage

Hour Maximum Average